ADVERSARIAL MACHINE LEARNING

ATTACK MODELS: CARLINI WAGNER L2 ATTACK

23

Carlini, N., & Wagner, D. (2017, May). Towards evaluating the robustness of neural networks.

- This is a hard problem. So they defined f such that:
- So the optimization problem becomes:
- Variable replacement
- So the optimization problem becomes
- They solve optimization using multiple starting point gradient descent

$\mathcal{C}(x+\delta) = \ell$ if and if only $f(x+\delta) \leq 0$

$$\text{minimize } \mathcal{D}(x, x+\delta) + c \cdot f(x+\delta) \quad \text{s.t. } x+\delta \in [0,1]^n$$

They replaced δ_i as $\delta_i = \frac{1}{2}(\tanh(w_i) + 1) - x_i$

$$\text{minimize} \quad \|\frac{1}{2}(\tanh(w)+1) - x\|_2^2 + c \cdot f(\frac{1}{2}(\tanh(w)+1)$$

with f defined as

$$f(x') = \max(\max\{\mathcal{H}(x')_i : i \neq l\} - \mathcal{H}(x')_l, -\kappa).$$

PRACTICAL BLACK BOX ATTACK

- Papernot et al. proposed a practical black box attack on CNN.
- They first train a substitute DNN on target classifier (oracle õ)
- They used Adversarial Examples crafted on substitute DNN to attack oracle.

Algorithm 1 - Substitute DNN Training: for oracle $\tilde{O}$, a maximum number max_ρ of substitute training epochs, a substitute architecture F, and an initial training set S_0.

Require: $\tilde{O}$, max_ρ, S_0, λ

1: Define architecture F
2: **for** $\rho \in 0 \,..\, max_\rho - 1$ **do**
3: *// Label the substitute training set*
4: $D \leftarrow \left\{(\vec{x}, \tilde{O}(\vec{x})) : \vec{x} \in S_\rho\right\}$
5: *// Train F on D to evaluate parameters θ_F*
6: $\theta_F \leftarrow \text{train}(F, D)$
7: *// Perform Jacobian-based dataset augmentation*
8: $S_{\rho+1} \leftarrow \{\vec{x} + \lambda \cdot \text{sign}(J_F[\tilde{O}(\vec{x})]) : \vec{x} \in S_\rho\} \cup S_\rho$
9: **end for**
10: **return** θ_F

Papernot, N., McDaniel, P., Goodfellow, I., Jha, S., Celik, Z. B., & Swami, A. (2017, April). Practical black-box attacks against machine learning.

ATTACK MODELS: CARLINI WAGNER L2 ATTACK

- This is a hard problem. So they defined f such that:

$$\mathcal{C}(x+\delta) = \ell \text{ if and if only } f(x+\delta) \leq 0$$

- So the optimization problem becomes:

$$\text{minimize } \mathcal{D}(x, x+\delta) + c \cdot f(x+\delta) \quad \text{s.t. } x+\delta \in [0,1]^n$$

- Variable replacement

$$\text{They replaced } \delta_i \text{ as } \delta_i = \tfrac{1}{2}(\tanh(w_i)+1) - x_i$$

- So the optimization problem becomes

$$\text{minimize } \|\frac{1}{2}(\tanh(w)+1) - x\|_2^2 + c \cdot f(\frac{1}{2}(\tanh(w)+1)$$

$$\text{with } f \text{ defined as}$$

$$f(x') = \max(\max\{\mathcal{H}(x')_i : i \neq l\} - \mathcal{H}(x')_l, -\kappa).$$

- They solve optimization using multiple starting point gradient descent

Carlini, N., & Wagner, D. (2017, May). Towards evaluating the robustness of neural networks.